


Sensor Fusion – 2D feature tracking - ReadMe

By Gaurav Borgaonkar

(25 Nov 2019)



The aim of this project is to implement and test different types of keypoint detector and keypoint descriptor algorithms and benchmark their performance in order to build a camera based collision detection system in the future.

The steps followed are as follows –

1. To begin with, we will be loading images, setting up data structures and putting everything into a ring buffer to optimize the memory load.
2. Then, several keypoint detectors such as HARRIS, FAST, BRISK and SIFT and compared with respect to number of keypoints and speed.
3. In the next part, we will then focus on descriptor extraction and matching using brute force approach and the FLANN approach.
4. In the last part, once the code framework is complete, we will test the algorithms in different combinations and compare them with respect to some performance measures.

Dependencies for Running Locally:

- `cmake` ≥ 2.8
 - All OSes: [click here for installation instructions] (<https://cmake.org/install/>)
- `make` ≥ 4.1 (Linux, Mac), 3.81 (Windows)
 - Linux: `make` is installed by default on most Linux distros
 - Mac: [install X-code command line tools to get `make`] (<https://developer.apple.com/xcode/features/>)
 - Windows: [Click here for installation instructions] (<http://gnuwin32.sourceforge.net/packages/make.htm>)
- OpenCV ≥ 4.1
 - This must be compiled from source using the `-D OPENCV_ENABLE_NONFREE=ON` `cmake` flag for testing the SIFT and SURF detectors.
 - The OpenCV 4.1.0 source code can be found [here] (<https://github.com/opencv/opencv/tree/4.1.0>)
- `gcc/g++` ≥ 5.4
 - Linux: `gcc` / `g++` is installed by default on most Linux distros
 - Mac: same deal as `make` - [install X-code command line tools] (<https://developer.apple.com/xcode/features/>)
 - Windows: recommend using [MinGW] (<http://www.mingw.org/>)

Basic Build Instructions

1. Clone this repo.
2. Make a build directory in the top level directory: ``mkdir build && cd build``
3. Compile: ``cmake .. && make``
4. Run it: ``./2D_feature_tracking``.

Project Rubric Points:

| Task No | Task | Implementation Description |
|---------|--|---|
| 1 | Data Buffer Optimization - Implement a vector for data buffer objects whose size does not exceed a limit (e.g. 2 elements). This can be achieved by pushing in new elements on one end and removing elements on the other end. | Implemented a basic "if" loop that checks if the size of the Data Buffer vector with 2. If the size is more than 2, it will remove the element at the beginning of the vector. (i.e. The oldest image loaded is removed from the data buffer.) |
| 2 | Keypoint Detection - Implement detectors HARRIS, FAST, BRISK, ORB, AKAZE, and SIFT and make them selectable by setting a string accordingly. | Created a feature detector object. When the input string matches, that detector type is assigned to the object and executed. Wrote two functions, one for Harris Corner Detection and another for modern detectors FAST, BRISK, ORB, AKAZE and SIFT |
| 3 | Keypoints removal - Remove all keypoints outside of a pre-defined rectangle and only use the keypoints within the rectangle for further processing. | I created a new keypoint vector for filtered keypoints. If the keypoint is located in the specified (mostly center) rectangular space in the image, we add it to filtered keypoints vector. |
| 4 | Keypoint Descriptors - Implement descriptors BRIEF, ORB, FREAK, AKAZE and SIFT and make them selectable by setting a string accordingly. | Created a descriptor extractor object. When the input string matches, that descriptor type is assigned to the object and executed. Wrote cases for BRISK, BRIEF, ORB, FREAK, AKAZE and SIFT descriptors |
| 5 | Descriptor Matching - Implement FLANN matching as well as k-nearest neighbor selection. Both methods must be selectable using the respective strings in the main function. | Created string variables to select descriptor matcher and implemented FLANN matching in matching2D_student.cpp file. Debugged the file for opencv bug by converting binary descriptors into floating point vectors |
| 6 | Descriptor Distance Ratio - Use the K-Nearest-Neighbor matching to implement the descriptor distance ratio test, which looks at the ratio of best vs. second-best match to decide whether to keep an associated pair of keypoints. | Used k = 2 and a threshold value of 0.8 between True Positives and False Positives for descriptor matching |

| | | |
|---|--|--|
| 7 | Your seventh task is to count the number of keypoints on the preceding vehicle for all 10 images and take note of the distribution of their neighborhood size. Do this for all the detectors you have implemented. | I created a variable total_keypoints that is updated over every instance of the for loop execution and gives us the total number of keypoints at the end of execution of the program. |
| 8 | Your eighth task is to count the number of matched keypoints for all 10 images using all possible combinations of detectors and descriptors. In the matching step, use the BF approach with the descriptor distance ratio set to 0.8. | I created a variable total_matched_keypoints that adds the number of keypoints matched in every single execution of the for loop (for every pair of images) and at the end of execution of the program, gives us total. |
| 9 | Your ninth task is to log the time it takes for keypoint detection and descriptor extraction. The results must be entered into a spreadsheet and based on this information you will then suggest the TOP3 detector / descriptor combinations as the best choice for our purpose of detecting keypoints on vehicles. Finally, in a short text, please justify your recommendation based on your observations and on the data you collected. | <p>Similar to counting the number of keypoints, I created a time variable total_time that counts the time required for each pair of images for both detection and extraction and adds them together. This time variable adds times for all pairs of images together.</p> <p>After adding the times, we divide the total number of keypoints by total time required and calculate the efficiency of the detector/descriptor combination.</p> |

TASK MP. 7: To count the total number of keypoints on the preceding vehicle for every detector

| Detector Type | Keypoints detected for the preceding vehicle |
|---------------|--|
| SHITOMASI | 1179 |
| HARRIS | 176 |
| FAST | 4094 |
| BRISK | 2762 |
| ORB | 1161 |
| AKAZE | 1670 |
| SIFT | 1386 |

Performance Evaluation:

Fastest three detector-descriptor combinations have been highlighted in green.

| Sr No | Detector + Descriptor | Total KeyPoints | Total Matches | Total Time (ms) | Ratio (matches/ms) |
|-------|-----------------------|-----------------|---------------|-----------------|--------------------|
| 1 | SHITOMASI + BRISK | 13423 | 4912 | 338.80 | 14.50 |
| 2 | SHITOMASI + BRIEF | 13423 | 6625 | 219.72 | 30.15 |
| 3 | SHITOMASI + ORB | 13423 | 5673 | 220.91 | 25.17 |
| 4 | SHITOMASI + FREAK | 13423 | 4084 | 622.21 | 6.564 |
| 5 | SHITOMASI + AKAZE | 13423 | NaN | NaN | NaN |
| 6 | SHITOMASI + SIFT | 13423 | 9484 | 474.33 | 19.99 |
| 7 | HARRIS + BRISK | 728 | 410 | 202.91 | 2.02 |
| 8 | HARRIS + BRIEF | 728 | 477 | 180.05 | 2.65 |
| 9 | HARRIS + ORB | 728 | 427 | 193.54 | 2.21 |
| 10 | HARRIS + FREAK | 728 | 290 | 543.42 | 0.53 |
| 11 | HARRIS + AKAZE | 728 | NaN | NaN | NaN |
| 12 | HARRIS + SIFT | 728 | 554 | 384.82 | 1.439 |
| 13 | FAST + BRISK | 49204 | 14267 | 503.38 | 28.34 |
| 14 | FAST + BRIEF | 49204 | 19554 | 195.83 | 99.85 |
| 15 | FAST + ORB | 49204 | 16896 | 107.25 | 157.54 |
| 16 | FAST + FREAK | 49204 | 11717 | 799.77 | 14.65 |
| 17 | FAST + AKAZE | 49204 | NaN | NaN | NaN |
| 18 | FAST + SIFT | 49204 | 30082 | 2134.49 | 14.09 |
| 19 | BRISK + BRISK | 27116 | 9563 | 4139.32 | 2.31 |
| 20 | BRISK + BRIEF | 27116 | 11426 | 3924.53 | 2.91 |
| 21 | BRISK + ORB | 27116 | 7039 | 3879.17 | 1.81 |

| | | | | | |
|----|---------------|-------|---------------------|---------------------|---------------------|
| 22 | BRISK + FREAK | 27116 | 7780 | 4472.89 | 1.74 |
| 23 | BRISK + AKAZE | 27116 | NaN | NaN | NaN |
| 24 | BRISK + SIFT | 27116 | 15279 | 6428.55 | 2.38 |
| 25 | ORB + BRISK | 5000 | 2427 | 141.03 | 17.21 |
| 26 | ORB + BRIEF | 5000 | 1822 | 109.10 | 16.70 |
| 27 | ORB + ORB | 5000 | 1996 | 134.08 | 14.89 |
| 28 | ORB + FREAK | 5000 | 920 | 522.93 | 1.76 |
| 29 | ORB + AKAZE | 5000 | NaN | NaN | NaN |
| 30 | ORB + SIFT | 5000 | 3360 | 1520.63 | 2.21 |
| 31 | AKAZE + BRISK | 13429 | 6869 | 1007.71 | 6.81 |
| 32 | AKAZE + BRIEF | 13429 | 7008 | 915.58 | 7.65 |
| 33 | AKAZE + ORB | 13429 | 5111 | 940.72 | 5.43 |
| 34 | AKAZE + FREAK | 13429 | 5319 | 1313.15 | 4.05 |
| 35 | AKAZE + AKAZE | 13429 | 8036 | 1654.58 | 4.86 |
| 36 | AKAZE + SIFT | 13429 | 9489 | 1478.94 | 6.42 |
| 37 | SIFT + BRISK | 13860 | 4285 | 1391.25 | 3.08 |
| 38 | SIFT + BRIEF | 13860 | 4553 | 1357.28 | 3.36 |
| 39 | SIFT + ORB | 13860 | Out of Memory Error | Out of Memory Error | Out of Memory Error |
| 40 | SIFT + FREAK | 13860 | 3334 | 2334.01 | 1.43 |
| 41 | SIFT + AKAZE | 13860 | NaN | NaN | NaN |
| 42 | SIFT + SIFT | 13860 | 7020 | 2298.95 | 3.05 |

Created a variable that adds the number of all keypoints after every iteration. I also created a time variable that adds together all the times for detector descriptor executions for every image pair sequence.

TASK MP. 7

For this, I created a variable **total_preceding_vehicle_keypoints** which updates itself for every loop execution of the image pair sequence. At the end, it prints the value of total number of preceding vehicle keypoints for the particular detector on screen.

TASK MP. 8

We count the number of matched keypoints for every pair of image sequence and store it in a variable **total_matches**. This variable is updated for every pair of images and at the end, it prints a value of total number of matched keypoints.

TASK MP. 9

Top Three Detector/Descriptor pairs:

The 3 best performing detector/descriptor pairs are selected based on their performance. Performance is evaluated by no. of matches on the preceding vehicle divided by time required.

Note: - For object tracking, although we are concerned about the keypoints only on the preceding vehicle; I have calculated the efficiency of these algorithms by using keypoints from all over the image.

| Sr No | Detector + Descriptor | Total KeyPoints | Total Matches | Total Time (ms) | Ratio (matches/time) |
|-------|-----------------------|-----------------|---------------|-----------------|----------------------|
| 1 | FAST + ORB | 49204 | 16896 | 107.25 | 157.54 |
| 2 | FAST+ BRIEF | 49204 | 19554 | 195.83 | 99.85 |
| 3 | SHITOMASI + ORB | 13423 | 6625 | 219.72 | 30.15 |